

COVASANT

Enterprise AI Platform

CAMS

Covasant Agent Management Suite

Product Documentation | Release V1.0 | 2026

The Operating System for Enterprise AI Agents

www.covasant.com

1. Executive Summary

Covasant Agent Management Suite (CAMS) is an enterprise-grade platform purpose-built to manage the complete lifecycle of AI agents: from strategic planning and development through orchestration, governance, and live production monitoring. CAMS addresses the full operational complexity of deploying AI at scale within enterprise environments.

Organizations deploying AI today face a fragmented landscape: strategy tools are disconnected from development environments, and production monitoring exists in isolation from build and deploy workflows. CAMS unifies all stages of the AI agent lifecycle into a single platform, eliminating tool sprawl and reducing time to production for enterprise AI initiatives.

Platform Vision: CAMS is not a point tool. It is the operating system for enterprise AI agents: enabling organizations to move from strategy to production with governance, quality, and observability built in at every stage.

2. Product Overview

2.1 Problem Statement

Enterprises adopting AI agents today face a deeply fragmented operational landscape. Agent development happens in silos without standardized governance, deployment is disconnected from build workflows, and production monitoring lacks centralized visibility and control. This fragmentation results in:

- High time-to-production for AI initiatives due to lack of a unified build-to-deploy pipeline.
- Agent sprawl: untracked, ungoverned agents proliferating across teams and environments.
- Absence of quality gates, leading to untested agents being deployed to production.
- No centralized cost visibility or budget control across LLM usage, models, and providers.
- Compliance and governance risk from uncontrolled agent behavior and missing audit trails.
- Executive misalignment: AI projects lacking formal business case documentation or ROI measurement.

CAMS resolves these challenges by providing a single, agentic management suite for the entire lifecycle, from strategy through production monitoring.

CAMS is designed as a platform-agnostic capability layer that seamlessly integrates across diverse technology ecosystems, including hyper-scalers, enterprise platforms, low code / no code platforms and proprietary environments. This ensures organisations retain flexibility of choice while avoiding vendor lock-in, enabling consistent governance, orchestration, and scalability of AI solutions regardless of the underlying stack.

2.2 What is CAMS?

CAMS (Covasant Agent Management Suite) is Covasant's enterprise-grade AI agent lifecycle management platform. It provides a structured, end-to-end environment for organizations to build intelligent agents, orchestrate multi-agent workflows, govern agent deployments, and monitor production operations, all within a single unified interface.

CAMS is built for enterprise scale, supporting multi-environment deployments, multi-tenant configurations, role-based access control, policy enforcement, cost management, and audit logging. It supports both no-code and pro-code agent development paradigms, making it accessible to business analysts, developers, and data scientists alike.

2.3 Objectives

CAMS is designed to achieve the following core objectives for enterprise organizations:

#	Objective	Description
1	Unify the AI Agent Lifecycle	Consolidate build, orchestration, governance, and monitoring into a single platform, eliminating tool sprawl.
2	Accelerate Time-to-Production	Enable teams to move from approved use case to deployed, production-certified agent in a structured, repeatable workflow.
3	Enforce Enterprise Governance	Provide centralized catalog management, multi-environment controls, audit logging, and policy enforcement across all agent deployments.

#	Objective	Description
4	Control Costs	Deliver granular cost attribution by agent, model, provider, and user: with configurable budget caps and spend alerts.
5	Enable Responsible AI	Embed guardrails, PII protection, compliance policy enforcement, and ethics & risk assessment throughout the platform lifecycle.

2.4 Target Users

CAMS serves the following enterprise personas, each with a primary module and defined responsibilities within the AI agent lifecycle:

User Role	Primary Module	Responsibility
AI Developer / Engineer	Agent Work Bench	Build, configure, and manage no-code and pro-code AI agents; manage knowledge bases, tools, and model connections.
Solutions Architect	Multi-Agent Orchestration	Design and deploy multi-agent pipelines using the visual workflow canvas; define sequential and parallel execution flows.
Governance Lead / Platform Admin	Agent Registry	Govern agent catalog, manage deployment plans, control environment promotion, and maintain compliance audit trails.
AI Operations Manager	Agent Control Tower	Monitor live agent performance, manage LLM costs, enforce guardrail policies, and investigate production incidents.

For **executive stakeholders** tasked with scaling AI adoption, CAMS provides a structured and governed pathway to accelerate innovation while mitigating risk. It establishes enterprise-grade guardrails for compliance, auditability, and accountability, enabling leadership to confidently expand AI initiatives without compromising on governance, security, or regulatory obligations.

3. Platform Overview

CAMS organizes the AI agent journey into four stages, each represented by a dedicated module. Together, these modules form a continuous, governed lifecycle: from building agents to production and beyond.

Stage	Module	User Role	Primary Action
Build	Agent Work Bench	AI Developer / Engineer	Build no-code or pro-code agents with model, tools, and knowledge base
Connect	Multi-Agent Orchestration	Solutions Architect	Wire individual agents into coordinated multi-step workflows on a visual canvas
Register	Agent Registry	Governance Lead / Platform Admin	Sync agents to org-wide catalog, manage deployment plans and environments
Monitor	Agent Control Tower	AI Ops Manager	Monitor live agent performance, cost, logs, guardrails, and policy compliance

CAMS is built on a modular architecture, where each component can be independently deployed, scaled, and managed. This modularity allows organisations to adopt capabilities incrementally, integrate with existing systems, and tailor deployments based on specific business and technical requirements. As an example, if you already have Google Gemini enterprise AI or Microsoft AI Foundry, you do not have to adopt Agent Work Bench to build agents, therefore Agent Control Tower can be deployed as an independent module.

3.1 Key Differentiators

- End-to-end lifecycle coverage from strategic planning to production monitoring on a single platform.
- Support for both no-code agents (businessRAGs users) and pro-code agents (custom framework engineers).
- Enterprise-grade governance including multi-tenant support, role-based access, and audit logging.
- Standardized MCP (Model Context Protocol) tooling for secure, extensible agent-to-service integrations.
- Real-time cost attribution with per-agent, per-model, and per-provider spend tracking.

4. Module Reference

MODULE
4.1
AWB

Agent Work Bench

Core Agent Development and Management Environment

Agent Work Bench is the primary agent development environment within CAMS. It provides complete infrastructure for technical teams to create, configure, and manage individual AI agents. The Workbench supports No-Code agents and Pro-Code agents within a single unified workspace, eliminating the need for multiple disconnected tools.

The Workbench is built on a secure-by-design principle, embedding security, compliance, and governance controls throughout the agent development lifecycle. This ensures that every AI agent is designed, tested, and deployed with built-in safeguards, reducing risk while accelerating time to value.

4.1 Agent Work Bench - Key Capabilities

Capability	Description
Dashboard	Provides an at-a-glance view of the total number of agents across the workspace, with a clear breakdown by agent type (No-Code, Pro-Code and their current status (Active, Idle, Pending).
Agents	Supports two agent types: No-Code agents built through a guided interface, Pro-Code agents developed programmatically using custom development frameworks. Within each agent, users can configure Instructions, Skills, and Knowledge Bases, attach tools, and create multiple sub-agents. Supports LLM Agent, Sequential Agent, and Parallel Agent patterns.
Playground	An integrated playground environment that allows users to interact with and validate an agent immediately after configuration. Users can select the relevant model within the Playground to test responses across different LLM backends before deployment.
Knowledge Bases (RAG)	Supports multiple providers for both unstructured and structured data including Vertex AI, OpenRAG, Azure AI, AWS Kendra, and others. Enables creation and management of RAG pipelines that provide agents with domain-specific, grounded responses.
Tools	Supports multi-tool integration to agents via MCP. Provides a catalog of callable tools that agents can invoke during execution, with pre-built and custom tools assignable per agent through standardized interfaces.
Multi-Agent Orchestration (MAO)	Visual canvas-based orchestration supporting zoom in and out, with the ability to add multiple instances of agents, including both Agents build using Agent WorkBench and Agents built externally using 3 rd party platforms. MAO supports orchestrating 3 rd party Agents using Agent to Agent (A2A) Protocol. MAO supports end-to-end testing that visualizes the full orchestration pipeline: Preprocessing, Planning, Execution, and the final response to the user.
Model Hub	A repository of AI models that is continuously updated as new models become available. Users can browse, configure, and connect models to agents, with support for multi-provider credential management.

Capability	Description
MCP Hub	Central hub for Model Context Protocol server integrations, enabling agents to call external services and APIs through secure, auditable, and standardized interfaces.
Enterprise Data Connectors	The platform supports seamless integration with enterprise data ecosystems, including data lakes, warehouses, and existing infrastructure. Through secure connectors and APIs, CAMS enables AI agents to access governed, high-quality data while maintaining alignment with enterprise data management and security policies.

USER JOURNEY

The developer receives a Predictive Maintenance agent brief, creates a No-Code AI Agent, configures instructions and skills, attaches a Knowledge Base of historical maintenance logs, and assigns a search tool via MCP. The user opens the Playground, selects Gemini 2.5 Pro, and validates responses before deployment. A solutions architect simultaneously wires five agents into a Sequential pipeline on the MAO canvas, tests the full execution plan visually, and publishes.

Multi-Agent Orchestration

Visual Workflow Composition for Multi-Step AI Pipelines

Multi-Agent Orchestration (MAO) provides a visual, drag-and-drop canvas for connecting individual AI agents into coordinated multi-step pipelines. MAO enables solutions architects to design and deploy complex multi-agent pipelines without custom engineering effort.

4.2 Multi-Agent Orchestration (MAO) - Key Capabilities

Capability	Description
Visual Canvas	Drag-and-drop, node-based canvas with zoom in and out support. Place agent nodes, draw connections, define branching logic, and create parallel or sequential execution paths.
Multi-Instance Agent Support	Add multiple instances of agents to the canvas, including both Agents build using Agent WorkBench and Agents built externally using 3 rd party platforms enabling flexible and hybrid agentic orchestration usecases. MAO supports orchestrating 3 rd party Agents using Agent to Agent (A2A) Protocol.
Sequential & Parallel Paths	Supports sequential agent chains and parallel branch-and-merge patterns, covering complex enterprise workflow topologies.
Playground	MAO has a in-built playground to quickly test the orchestration flow and to visualize the full processing pipeline: Preprocessing, Planning Executor, and Response Synthesis, enabling quick validation before promoting it to QA, Staging or production deployments.
MAO Library	Searchable catalog of all orchestrations with name, description, status (Active/Draft), instruction count, instance count, and last-updated metadata.
Draft & Publish Controls	Save orchestrations as Draft during design and testing, then Publish to promote it to QA, Staging or production deployments..
Embedded Workbench Access	MAO canvas accessible directly within Agent Work Bench, enabling developers to build and manage orchestrations without switching platform contexts.

USER JOURNEY

The solutions architect opens MAO and creates a Manufacturing Maintenance Pipeline. On the visual canvas, the user places the System Health Agent as the entry node, connects it to the Condition Agent for anomaly detection, branches to the Parts Agent and Workforce Agent running in parallel, converges them into the Confirmation Agent, and routes output to the Dashboard Agent. The orchestration is saved as Draft, tested with simulated sensor input, and published. The end-to-end predictive maintenance workflow is live.

Agent Registry

Organization-Wide Agent Catalog and Governance Layer

Agent Registry is the centralized governance and catalog layer for all AI agents deployed across the organization. It serves as the authoritative source of truth for what agents exist, where they are deployed, and how they are managed across environments and tenants.

4.3 Agent Registry - Key Capabilities

Capability	Description
Dashboard	Displays key stats including total number of agents, instances, and tenants, along with instance activity trends over the last period. Also shows a real-time activity log of all recent actions taken within the Registry.
Agent Library	Lists all registered agents, including agents built within CAMS as well as agents from external providers. Also shows the total instances of each agent across environments, with status and deployment details.
Providers	Manage and configure the 3 rd party Agent Platform providers so that Agent Registry can auto-discover your agents built and hosted on 3 rd Party providers in a single consolidated Agent Library. Agent Registry currently supports automatic discovery of Agents built on Gemini Enterprise, Azure AI Foundry, AWS Bedrock and Databricks AgentBricks.
Sync Agents	One-click synchronization keeps the Registry current with agents built or modified in Agent Work Bench.
Environments	Users can create and manage deployment environments such as Dev, QA, Test, and Production. Users can also define their promotion workflows between these environments. Environments can be configured to use different runtimes for host the agents. Agents can be deployed on Kubernetes platforms including GKE, AKS, and EKS, or serverless runtimes like GCP CloudRun or Azure AppFunctions. External Agent Runtimes like GCP Agent Engine and Azure AI Foundry are also supported.
Tenants	Users can define multiple tenants if Agents support Multi-Tenancy. Tenant specific behaviour of the Agent (e.g. prompt, knowledge bases used, secrets etc) can be customized per tenant.
Configuration Management	Agent Behaviour (applicable only for Agents built using CAMS Agent Workbench) can be customized before deploying the Agent to different environments, for example if the Agent is expected to use a different Knowledge Base (or datastore) in QA Environment than in Production etc. For Agents that support multi-tenancy, tenant specific configuration can be created and managed from the configuration management.
Deployment Management	Allows users to define and plan the Agent deployment, planning deployment involves selecting the Agent to be deployed, the runtime environment to which the Agent needs to be deployed, if the Agent supports multi-tenancy, selecting the set of tenants and tenant specific configuration to be used by the agent, all of this can be done using a visual canvas in Agent Registry's Deployment Plan.

USER JOURNEY

The Governance Lead opens Agent Registry after the team has built 20+ agents. The user filters the Agent Library to Active status and spots an agent still showing as Pending. The user reviews the configuration, identifies a missing provider credential, and routes it to engineering. The user checks the Activity Log to confirm recent deployments completed without errors and reviews the Dashboard for a full summary of agents, instances, and tenant activity.

Agent Control Tower is the production operations center for all live AI agents running within the CAMS platform and 3rd party agents that use Agent Control Tower as an AI Gateway. It provides end-to-end visibility into agent activity, infrastructure costs, performance metrics, guardrail enforcement, LLM call logs, and policy compliance. Control Tower gives operations teams the real-time intelligence and control needed to run enterprise AI safely and cost-effectively at scale.

4.4 Agent Control Tower - Key Capabilities

Capability	Description
Dashboard	Provides a consolidated view of recent spending activity, platform usage metrics, and model-level activity tracking across all models in use.
Cost Attribution	Ranked views of highest-cost agents, models, users, and providers. Per-request cost tracking enables precise budget optimization and anomaly detection.
LLM Logs	Paginated, searchable log of every LLM API call: timestamp, status, session ID, agent name, model, provider, tokens, and per-call cost.
Policy Enforcement: Guardrails	Configure and enforce guardrail policies: content filters, PII detection and redaction, output length limits, and compliance policies (GDPR, HIPAA).
Agents View	Live operational status of all deployed agents with success rates, provider information, and Active/Inactive status for rapid anomaly detection.
Model Hub	Manage AI models available within the platform: register new models, configure parameters, and review usage statistics.
Providers Management	Configure external AI service providers (GCP, AWS, Azure, Salesforce) with API keys and service endpoints.
Budget Management	Set platform-wide or per-agent maximum budget thresholds to prevent cost overruns.
Evaluations	Enables AI Operations Managers to periodically run evaluations on the Agent to access the Agent's Performance, Accuracy and other key metrics. Agent Control Tower supports GCP Vertex AI Eval SDK and Weights and Biases Evaluation frameworks to run evaluations on the Agents.

USER JOURNEY

The AI Operations Manager opens Control Tower and reviews the weekly metrics dashboard: request volumes, success rates, and total spend are all within expected thresholds. In LLM Logs, the user spots a cluster of high-token-usage requests from the MAO pipeline making more orchestration hops than expected. The user configures a max token limit per agent call in Guardrails and sets a monthly platform budget cap. Agents are running, costs are controlled, and guardrails are enforced.

5. Technical Architecture

5.1 Supported Agent Types

CAMS supports two distinct agent development paradigms, enabling teams to choose the approach that best matches their technical capabilities and use case complexity:

Agent Type	Description
No-Code AI Agents	Configured through CAMS's guided interface. Suitable for business users and developers who need to deploy intelligent agents quickly without custom programming. Supports persona configuration, knowledge base attachment, tool assignment, and model selection through a graphical interface.
Pro-Code AI Agents	Built programmatically using custom development frameworks. Suitable for engineers requiring fine-grained control over agent behavior, sub-agent composition, and custom execution logic. Integrated into the CAMS platform via the Pro-Code Agent registration workflow.

5.2 Integration Standards

- Model Context Protocol (MCP): All external tool integrations are implemented as standardized MCP servers, enabling secure, auditable, and extensible agent-to-service communication.
- LLM Provider Support: Supports multiple LLM providers. Currently configured with Gemini 2.5 Pro via Google Vertex AI, with extensible provider management for additional LLM backends.
- RAG Knowledge Bases: Agents can be equipped with domain-specific Retrieval-Augmented Generation knowledge bases for context-grounded responses.

5.3 Deployment Architecture

- Multi-Environment Support: Agents are promoted through discrete environments (development, staging, production) with promotion controls managed through Agent Registry.
- Multi-Tenant Configuration: Enterprise-grade tenant isolation ensures organizational separation of agent catalogs, instances, and operational data.
- Audit and Compliance Logging: All platform actions are captured in structured activity logs with timestamps, actor attribution, and action categorization.
- Role-Based Access Control: Granular permission management across all platform modules, configurable at workspace and organization level.

5.4 Supported LLM Providers & Evaluation Frameworks

Category	Supported Options
LLM Providers	Google Vertex AI (Gemini 2.5 Pro: primary production model); additional providers configurable via Models and Providers management
Agent Frameworks	No-Code (guided CAMS interface); Pro-Code (custom development frameworks for stateful multi-agent workflows)

Category	Supported Options
Integration Protocol	Model Context Protocol (MCP): standardized interface for all external tool and service integrations. A2A Protocol: For Agent to Agent interface, primarily used as part of Multi-Agent Orchestration.

6. Enterprise Governance and Compliance

CAMS is designed with enterprise governance as a foundational requirement, not an afterthought. Governance controls are embedded across all modules, ensuring that AI agents are built, deployed, and operated within clearly defined organizational policies and compliance boundaries.

6.1 Governance Controls by Module

Module	Governance Controls
Agent Work Bench	Workspace-level isolation, role-based access control, agent status management (Active/Pending review states)
Agent Registry	Centralized catalog with full deployment lifecycle tracking, multi-tenant environment isolation, provider management, real-time activity audit log
Agent Control Tower	Real-time guardrail policy enforcement, budget and token limit controls, complete LLM call audit log, cost attribution by agent, model, and user

6.2 Responsible AI

CAMS incorporates responsible AI practices across the platform lifecycle. Agent Control Tower provides guardrail policy enforcement to constrain agent behaviour within approved content and operational boundaries.

7. Getting Started

7.1 Access & Authentication

CAMS is a web-based platform accessible via any modern browser (Chrome, Firefox, Edge, Safari: latest versions recommended). Users authenticate once using their organization's credentials via the KeyCloak authentication system and can seamlessly move between all CAMS applications without signing in again.

7.2 Recommended Workflow for New Users

Step	Action	CAMS Module
1	Build your agent: configure persona, select a base model, define system prompt, and optionally attach a knowledge base.	Agent Work Bench
2	Register and deploy: create a deployment plan, select target environment, and deploy the agent instance.	Agent Registry
3	Monitor in production: review dashboard metrics, monitor live sessions, investigate LLM logs, and adjust guardrails as needed.	Agent Control Tower

Note: Access credentials are provided by your organization's platform administrator. Contact your administrator to request an account and the appropriate access roles for your role.

8. Quick Reference

8.1 Module Summary

Module	Core Capabilities
Agent Work Bench	No-code and pro-code agent development, knowledge base management, tool assignment, model configuration, workspace and access management
Multi-Agent Orchestration	Visual drag-and-drop workflow canvas, parallel and sequential agent pipelines, draft/publish lifecycle, MAO library management
Agent Registry	Organization-wide agent catalog, multi-environment deployment management, provider configuration, audit activity log
Agent Control Tower	Real-time production monitoring, cost tracking by agent, model, provider, and user, LLM call audit log, guardrail policy enforcement, budget management

9. Glossary

This glossary defines all abbreviations, acronyms, and key terms used throughout the CAMS product documentation.

9.1 Abbreviations & Acronyms

Abbreviation	Full Form	Context
CAMS	Covasant Agent Management Suite	The end-to-end enterprise AI agent lifecycle management platform by Covasant
AI	Artificial Intelligence	The simulation of human intelligence by machines; foundational technology behind all CAMS capabilities
LLM	Large Language Model	The underlying AI model (e.g., Gemini, GPT-4, Claude, Llama) that powers agent reasoning and responses
RAG	Retrieval-Augmented Generation	A technique for grounding agent responses in domain-specific knowledge base content before generating an answer
MAO	Multi-Agent Orchestration	The CAMS module for composing individual agents into coordinated multi-step pipelines (Module 2)
REG	Agent Registry	Short reference for the Agent Registry module (Module 3)
CTR	Agent Control Tower	Short reference for the Agent Control Tower module (Module 4)
GCP	Google Cloud Platform	A supported external AI service provider configurable in Agent Control Tower
AWS	Amazon Web Services	A supported external AI service provider configurable in Agent Control Tower
ML	Machine Learning	A subset of AI; the discipline underpinning LLM training and evaluation methodologies
SDK	Software Development Kit	A set of tools used by developers to build pro-code agents and integrate with CAMS APIs
SLA	Service Level Agreement	Performance and availability commitments for deployed agent instances in production environments
SSO	Single Sign-On	Authentication mechanism; CAMS uses KeyCloak-based SSO so users authenticate once across all modules
A2A	Agent-to-Agent	Communication protocol enabling agents to invoke and coordinate with other agents as sub-agents within a workflow
GKE	Google Kubernetes Engine	A supported Kubernetes platform for deploying agent environments in Agent Registry

Abbreviation	Full Form	Context
AKS	Azure Kubernetes Service	A supported Kubernetes platform for deploying agent environments in Agent Registry
EKS	Elastic Kubernetes Service	A supported Kubernetes platform (AWS) for deploying agent environments in Agent Registry

9.2 Key Terms

Term	Definition
AI Agent	A software entity powered by an LLM, configured to perform specific tasks autonomously using instructions, tools, and knowledge bases
No-Code Agent	An AI agent configured through the CAMS graphical interface without custom programming; accessible to business users and developers
Pro-Code Agent	An AI agent built programmatically using custom development frameworks; suitable for engineers requiring fine-grained control
Knowledge Base	A collection of documents indexed and made searchable for RAG; attached to agents to provide domain-specific context
Guardrails	Policy enforcement mechanisms in Agent Control Tower that constrain agent behavior within approved content and operational boundaries
Tenant	A customer or organisational unit that is a user for the Agents built or managed through CAMS.
Instance	A running deployment of an agent in a specific environment (development, staging, or production)
Deployment Plan	A specification describing which agent to deploy, to which environment, with what infrastructure configuration
Session	A single conversation or task execution between a user or system and an agent
Trace	A detailed log of every step an agent took to produce a response, including all tool calls, latency, and errors
LLM Provider	An external LLM service (e.g., OpenAI, Anthropic, Google Vertex AI, AWS Bedrock) configured in CAMS for agent use
Tool	A callable capability (e.g., web search, code execution, API call) that an agent can invoke during task execution
Workspace	An organizational container in Agent Work Bench for grouping related agents by project, team, or function
Activity Log	A real-time, paginated audit trail of all platform actions with timestamps, actor attribution, and completion status
Orchestration	The coordination of multiple AI agents in sequence or parallel to complete complex, multi-step enterprise workflows